



CSCI-GA.3033-015

# Virtual Machines: Concepts & Applications

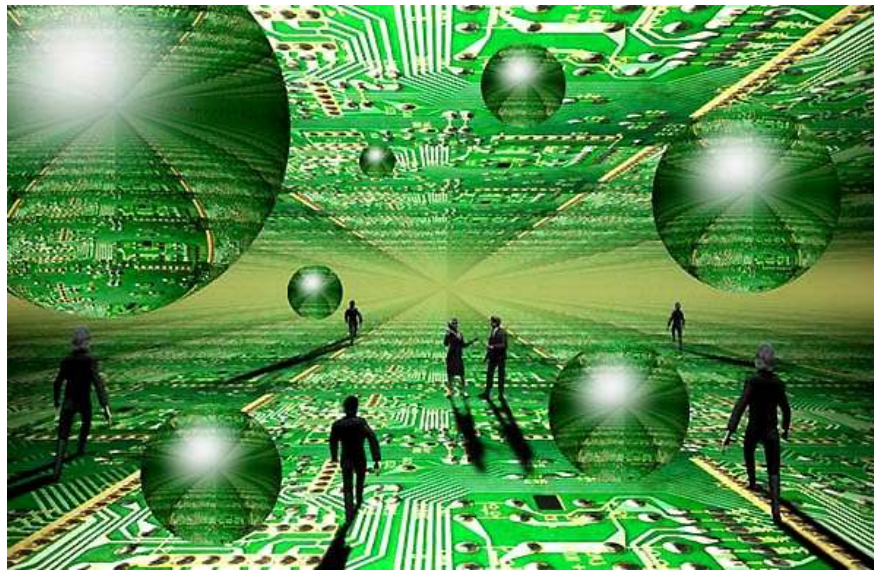
## Lecture 6: HLL VM - I

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>

**Disclaimer:** Many slides of this lecture are based on the slides of authors of the textbook from Elsevier.  
All copyrights reserved.



# What are we talking about?

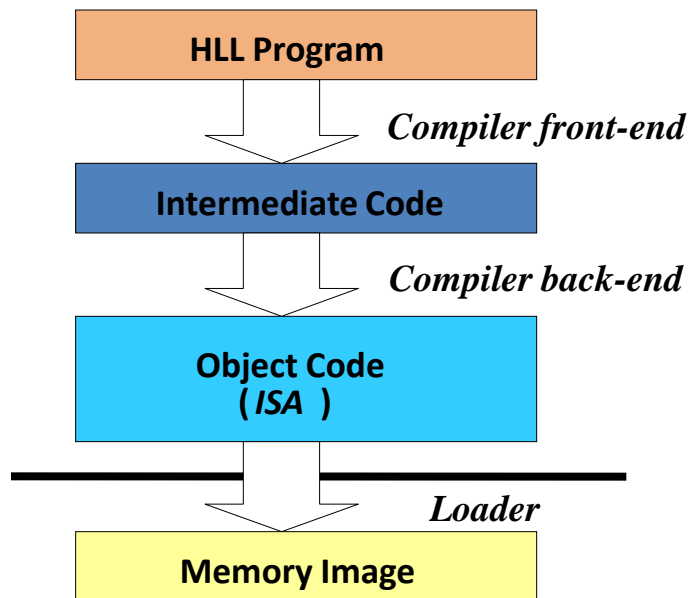
- Traditional process VMs are afterthoughts
  - How about if we *design* a special guest ISA/system interface:
    - With portability as the main goal
    - define an abstract interface that can be supported by all conventional OSes.
    - Reflects important features of specific HLL or class of HLLs.
    - Simplifies compilation
- ← This is why we call them HLL VMs

# HLL VM is similar to Process VM BUT ...

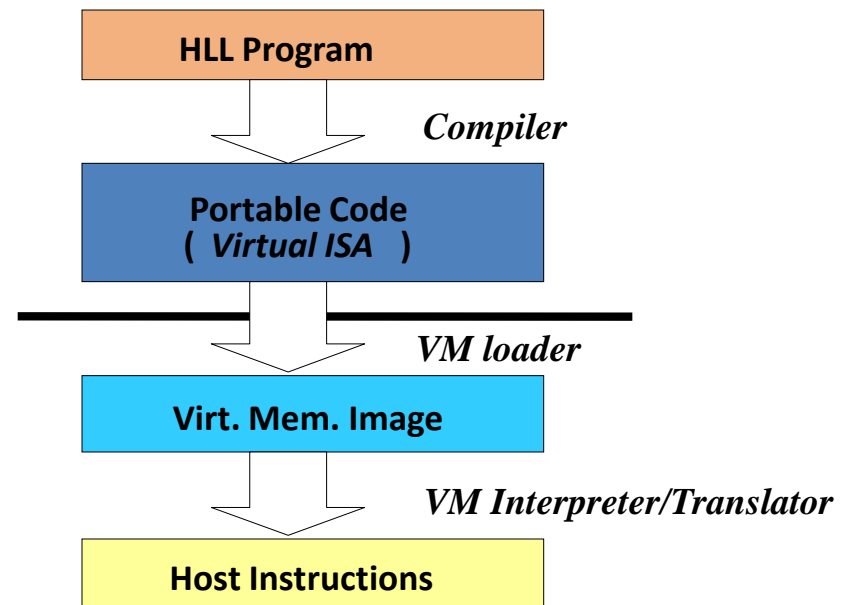
- ISA defined for user-mode programs only
- ISA not designed for real hardware
  - Only to be executed on virtual processor
  - Referred to as virtual-ISA or V-ISA
- System interface is a set of standardized APIs

# HLL VMs from language/compiler perspective

- Goal: complete **platform independence** for applications
- Virtual instruction set + libraries
  - Instead of ISA and OS interface



Traditional



HLL VM

# P-Code VM

- Popularized HLL VMs
- Provided highly portable version of Pascal
- Consists of
  - Primitive libraries
  - Machine-independent object file format
  - A set of byte-oriented “pseudo-codes”
  - Virtual machine definition of pseudo-code semantics

# P-Code VM

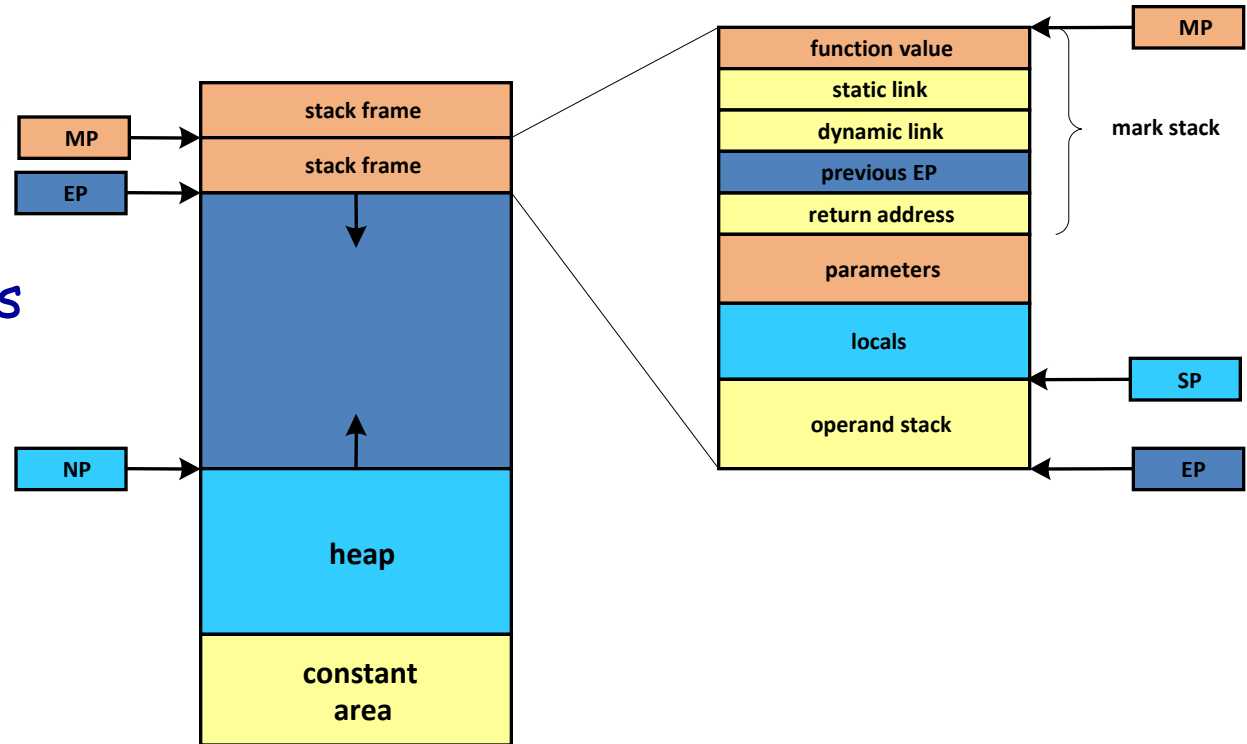
- Instruction set
  - Stack oriented
  - Stack "Frame" is part of VM definition

MP: Mark Pointer

EP: Extreme Pointer

NP: New Pointer

SP: Stack Pointer



```
lodi    0  3    // load variable from current frame (nest 0 depth),
              // offset 3 from top of mark stack.
ldci    1        // push constant 1
addi                    // add
stri    0  3    // store variable back to location 3 of current frame
```

# P-Code VM

- Advantages
  - Porting is simplified
    - Don't have to develop compilers for all platforms
  - VM implementation is smaller/simpler than a compiler
  - VM provides concise definition of semantics
  - Through interpretation, startup time is reduced
  - Generic I/O and Memory interface
    - Tended to be least common denominator
    - $\Rightarrow$  Relatively weak I/O capabilities

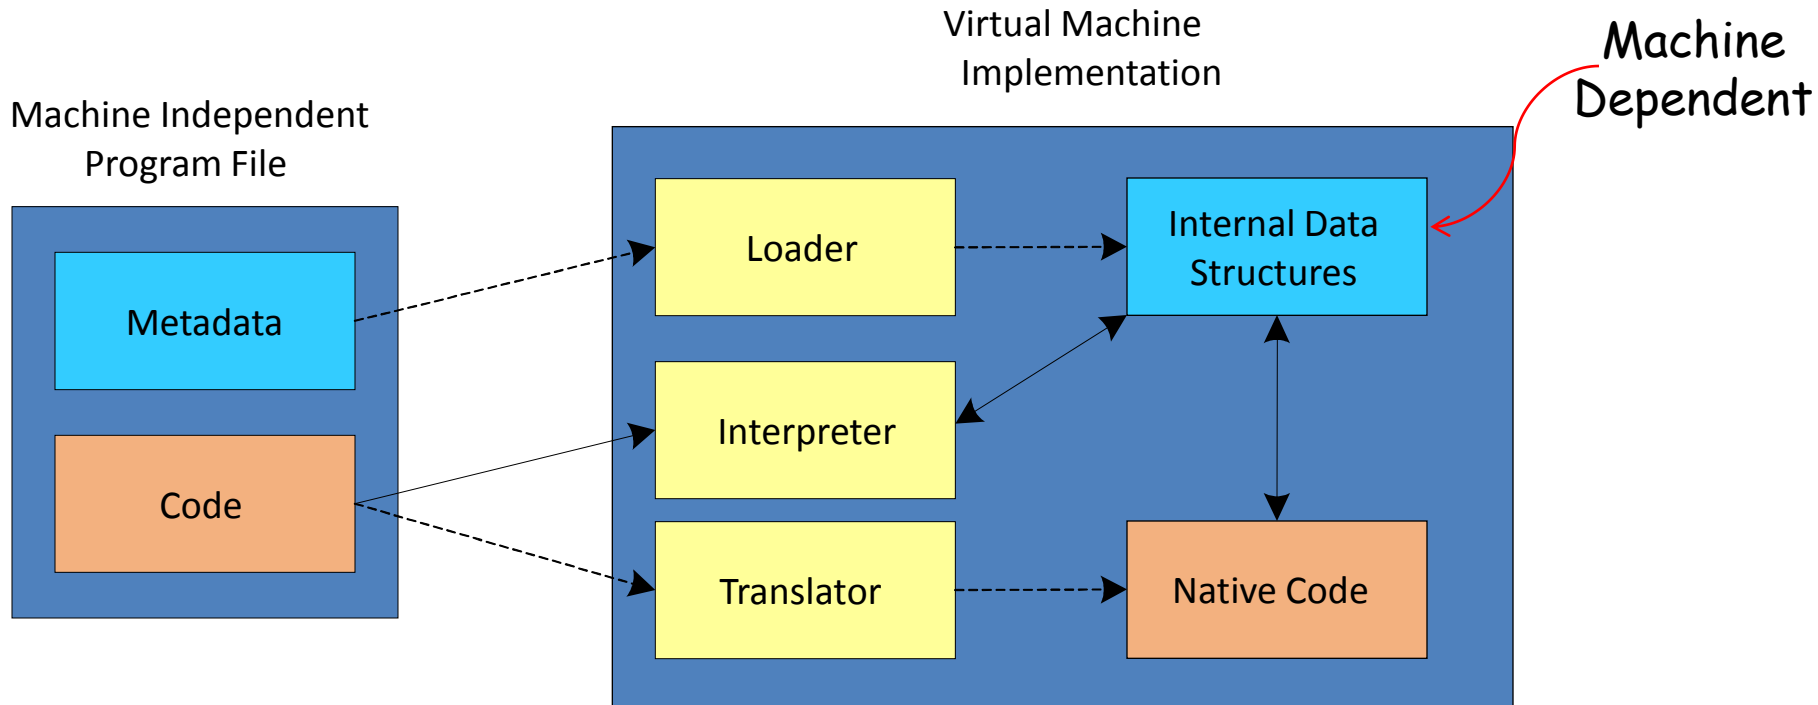
# Modern HLL VMs

- Superficially similar to P-code scheme
  - Stack-oriented ISA
  - Standard libraries
- Network Computing Environment
  - Untrusted software (this is the internet, after all)
  - Robustness (generally a good idea)
    - $\Rightarrow$  object-oriented programming
  - Bandwidth is a consideration
  - Good performance must be maintained
- Two major examples
  - Java VM
  - Microsoft Common Language Infrastructure (CLI)



# Modern HLL VMs

- Compiler forms program files (e.g. class files)
  - Standard format
- Program files contain both code and metadata



# Terminology

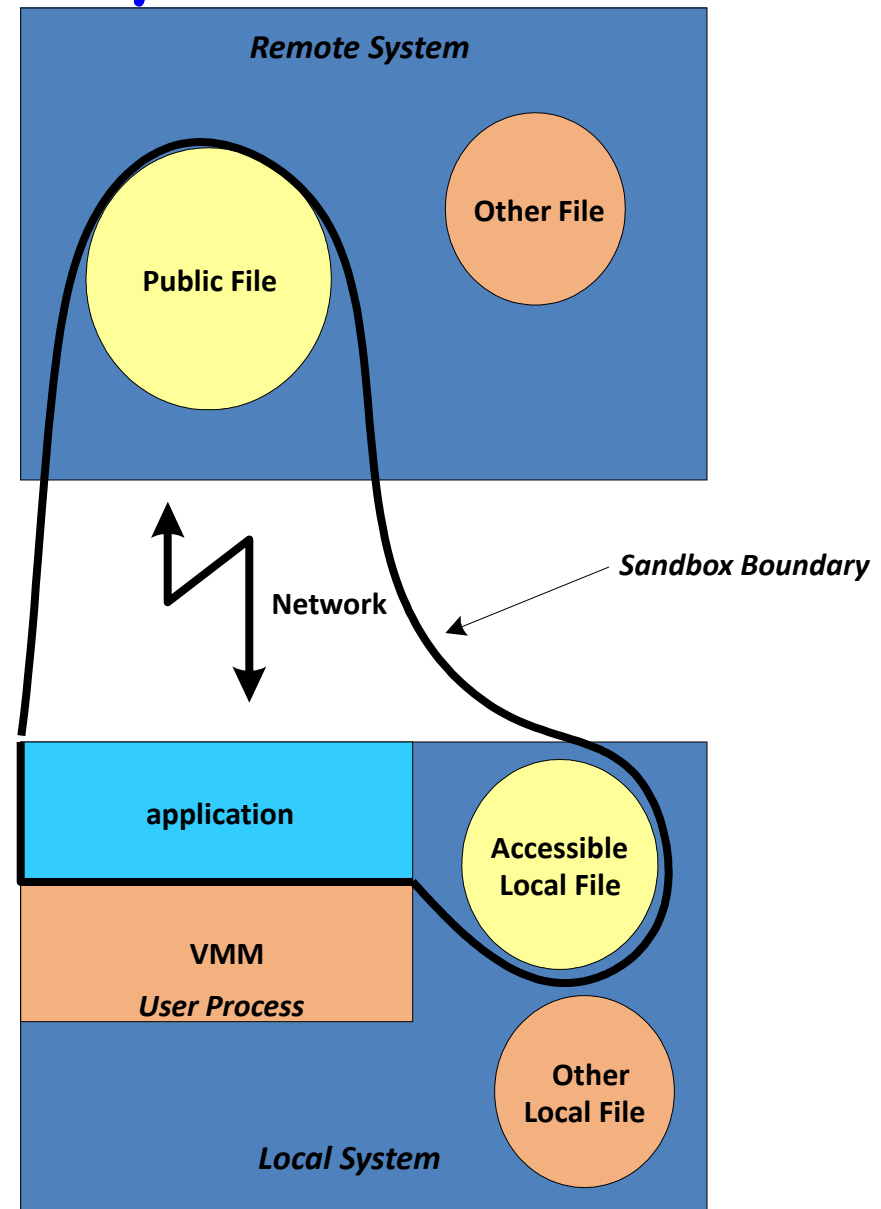
- Java Virtual Machine Architecture  $\Leftrightarrow$  CLI
  - Analogous to an ISA
- Java Virtual Machine Implementation  $\Leftrightarrow$  CLR (Common Language Runtime)
  - Analogous to a computer implementation
- Java bytecodes  $\Leftrightarrow$  Microsoft Intermediate Language (MSIL), CIL, IL
  - The instruction part of the ISA
- Java Platform  $\Leftrightarrow$  .NET framework
  - ISA + Libraries; a higher level ABI

# 4 Characteristics of HLL VMs

- Security
- Robustness
- Networking
- Performance

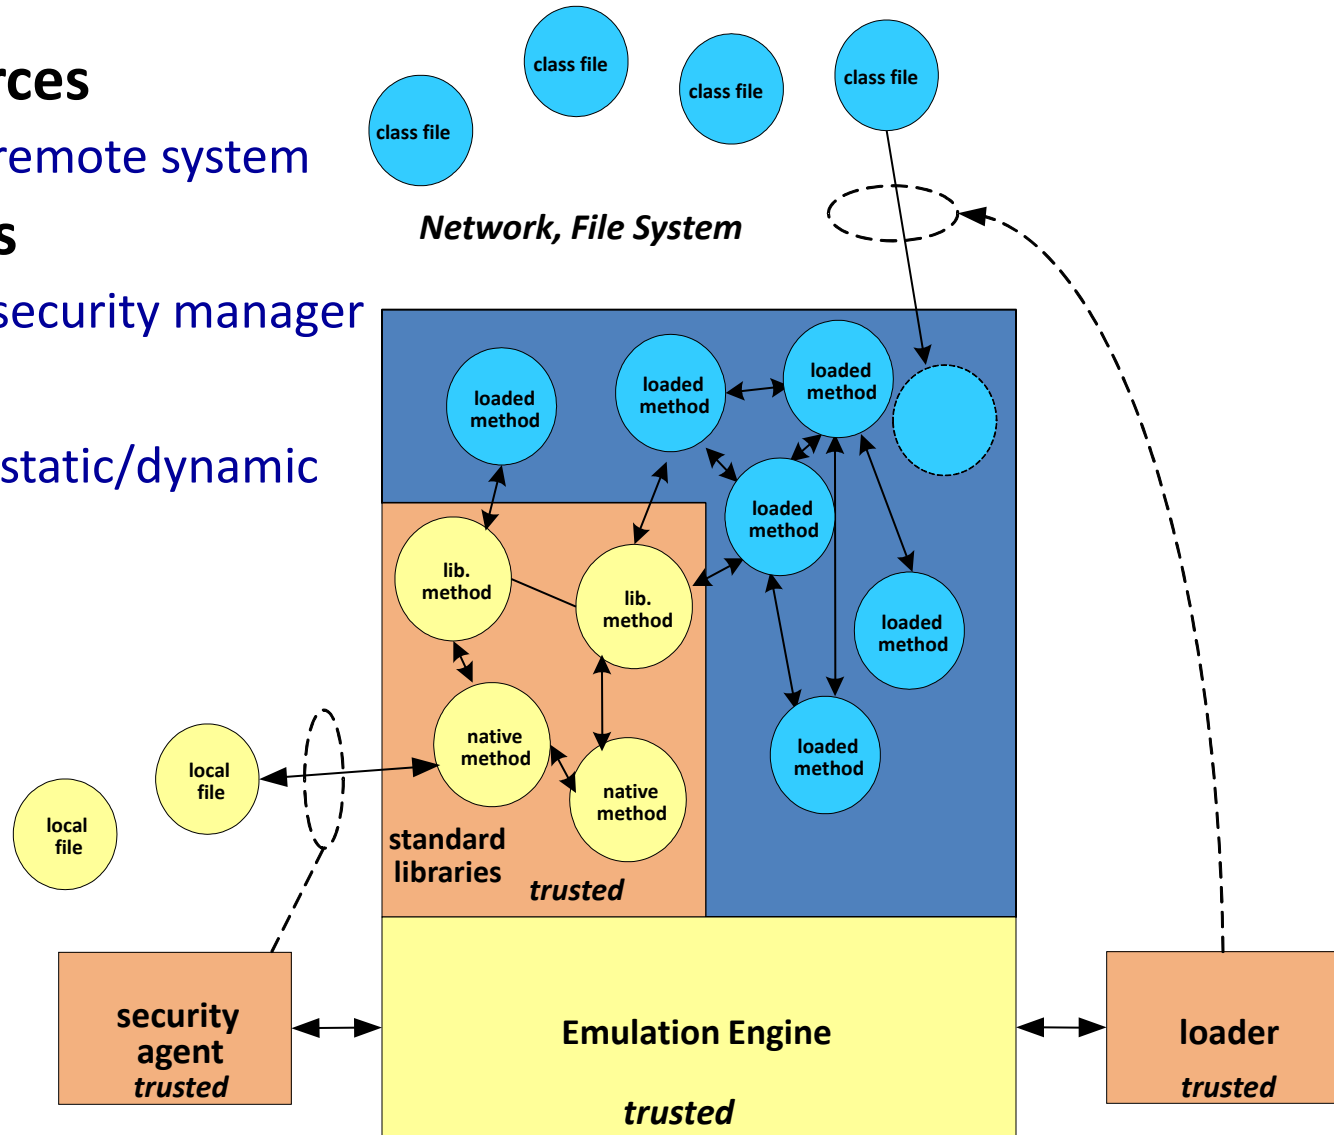
# Security

- A key aspect of modern network-oriented VMs
- Must protect:
  - Local files and resources
  - Runtime from user process
- The program runs in a **sandbox** at the host machine. It is **managed** by the VM runtime.
- The ability to load an **untrusted application** and run it in a managed secure fashion is a very big challenge!



# Protection Sandbox

- ❑ Remote resources
  - Protected by remote system
- ❑ Local resources
  - Protected by security manager
- ❑ VM software
  - Protected via static/dynamic checking



# Robustness: Object-Orientation

- Objects
  - Data carrying entities
  - Dynamically allocated
  - Must be accessed via pointers or *references*
- Methods
  - Procedures that operate on objects
- Class
  - A *type* of object and its associated methods
  - Object created at runtime is an *instance* of the class
  - Data associated with a class may be *dynamic* or *static*

OO programming paradigm has become the model of choice  
for modern HLL VMs.

Both Java and CLI are designed to support OO software.

# Networking

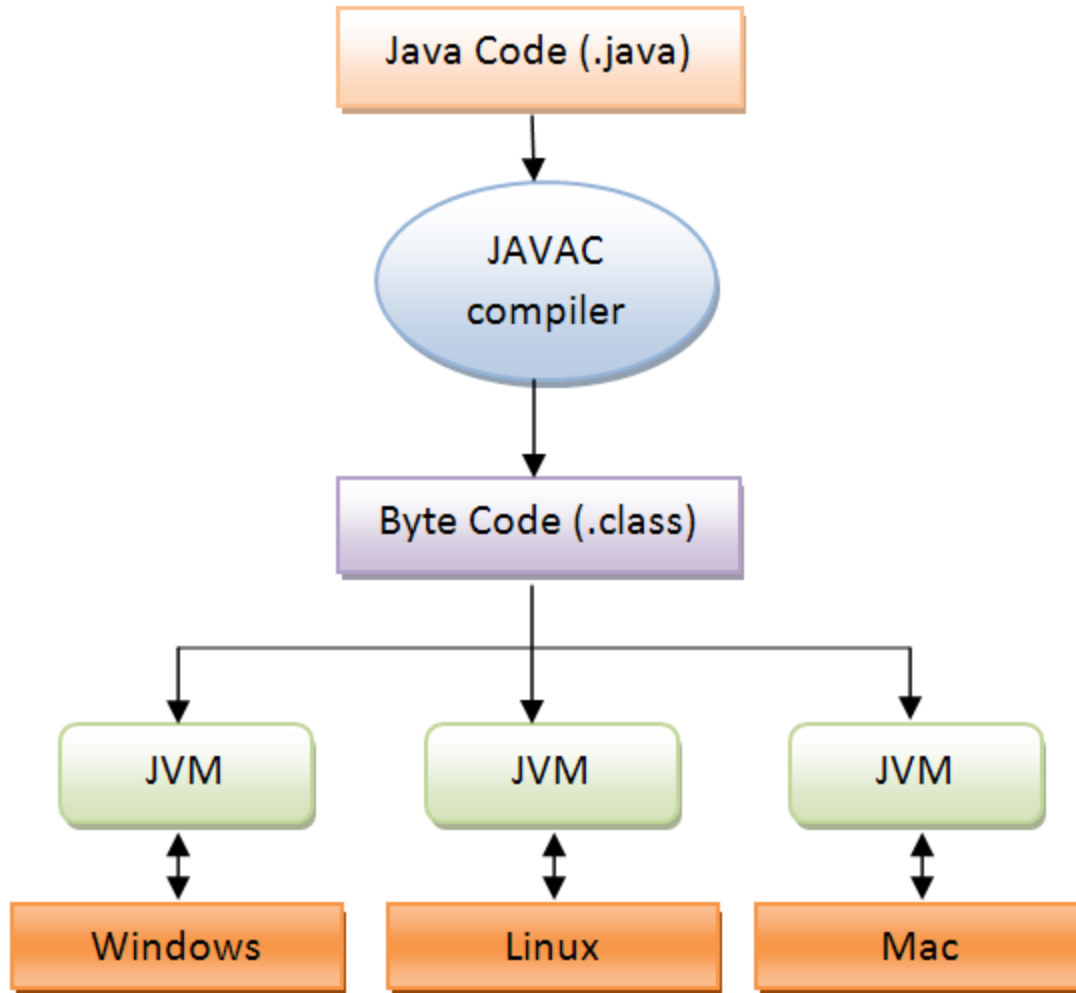
- The application must use the available bandwidth (scarce) efficiently
  - Application loaded incrementally → dynamic linking
  - Improves program startup-time

# Performance

- Of course we sacrifice some performance for the sake of portability
- Yet, we can use the techniques we learned so far (and some more as we proceed) to ensure good performance.



# Java Virtual Machine (JVM)



Source: <http://i.stack.imgur.com/Deo2s.png>

# Data items

- Types are defined, but not implementation details
  - Reference types (pointers): number of bits needed is not part of Java ISA.
  - Primitive types, e.g. int, char, byte, short, long, float, double,
  - Another primitive type: ReturnAddress (not in Java HLL but in Java ISA)
- Exact sizes of data types are not given
  - Only the range of values that can be held
    - e.g. byte is between -128 and +127

# Objects and Arrays

- Objects:
  - Logical structure, defined by programmer, to carry data
  - Composed of primitive data types and references
- Array
  - Fixed number of elements
  - All elements must be of the same type
  - If the elements are references then they must all point to objects of the same type

# Data Storage Types

- **Global**
  - the main memory
  - where globally declared variables reside
- **Local**
  - temporary storage
  - for variables local to a method
- **Operand**
  - holds variables while they are being operated on by functional instructions

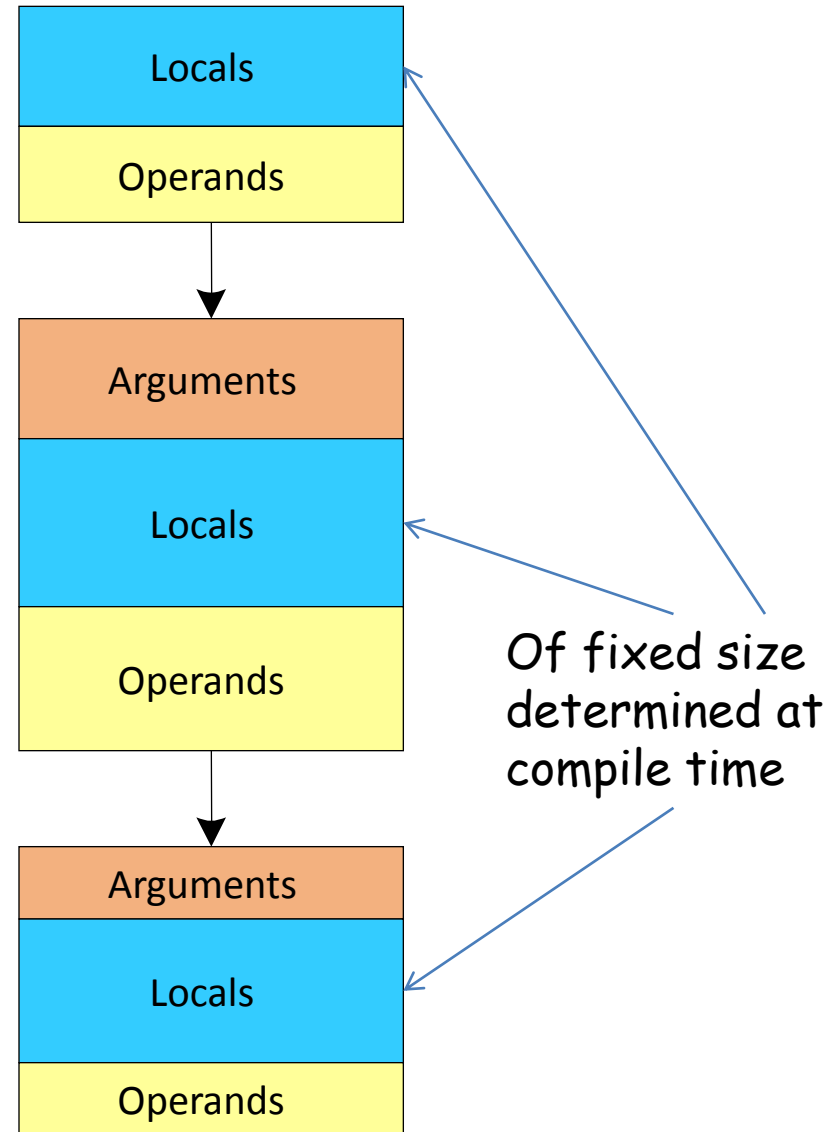
Allocated  
on the  
**stack**

# Data Storage Types

- All storage is divided into **cells** or **slots**
- A cell/slot usually holds a single data item
- Actual amount of bits needed for cell/slot is implementation dependent

# Stack

- Arguments
  - Locals
  - Operands
  - As each method is called, a stack frame is allocated.
- In that order



# Global Memory

- Method area
  - for holding code
- Global storage area
  - for holding arrays and objects
  - managed as a **heap**
  - of unspecified size with respect to JVM architecture
  - Can contain both static and dynamic objects

# Heap

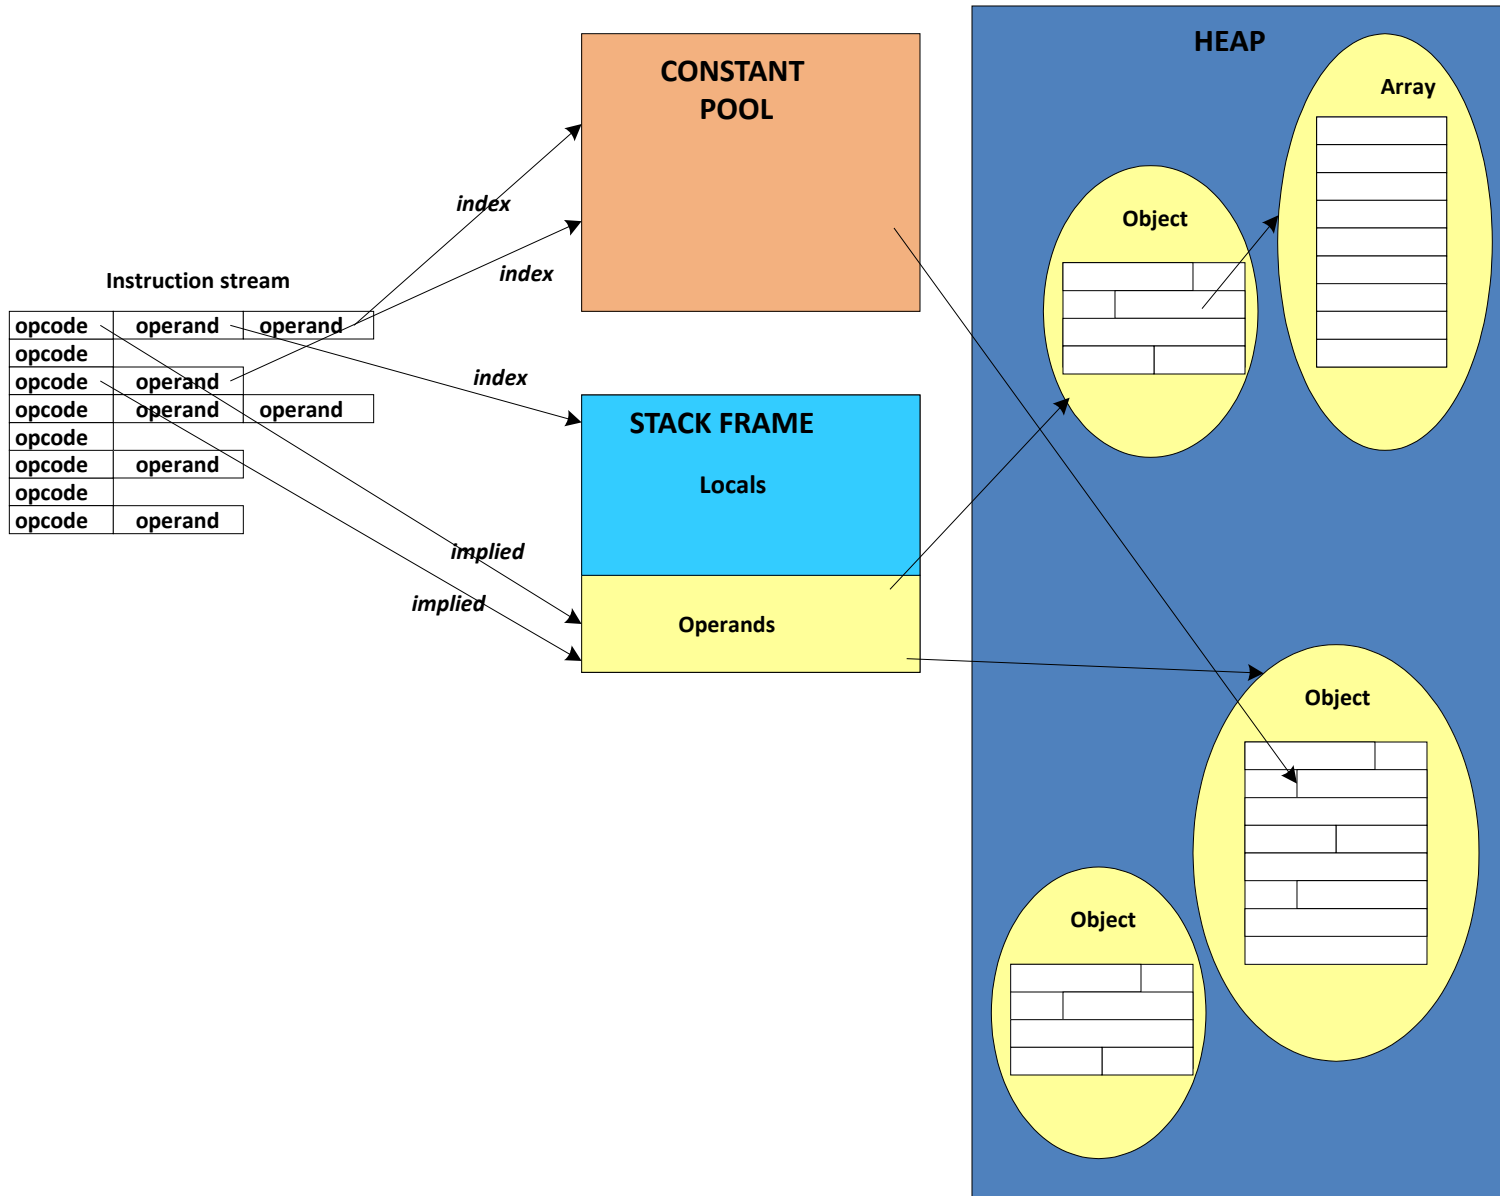
- Objects are created on heap
- Each application has only one
- Each application has its own
- JVM instructions allocate objects on heap
  - No instruction to release memory
  - Garbage collection is part of implementation
- Object representation is implementation dependent



# Constant Pool

- ISA allows constants to be expressed in the instruction as **immediate operands**
- But some constants:
  - are used by several instructions
  - are of different ranges
- So: Constant data associated with a program is placed in a block called constant pool
- Instructions access them by **indexing** constant pool
- Constant pool then:
  - **defined as part of the ISA**
  - Exact size of constants is specified
  - Does not change with program execution

# Putting is All Together: Memory Hierarchy in JVM



# Network Friendliness


- Support dynamic class file loading on demand
  - Load only classes that are needed
  - Spread loading out over time
- Compact instruction encoding
  - Use stack-oriented ISA (as in Pascal)

# Garbage Collected Heap


- Objects are created and “float” in memory space
  - Tethered by references
  - In architecture, memory is unbounded in size
  - In reality it is limited
- Garbage creation
  - During program execution, many objects are created then abandoned (become garbage)
- Collection
  - Due to limited memory space, Garbage should be collected so memory can be re-used
  - Forcing programmer to explicitly free objects places more burden on programmer
    - Can lead to memory leaks, reducing robustness
  - To improve robustness, have VM collect garbage automatically

# Instruction Set

- Stack based
- Defined for class file, not memory image
- Bytecodes
  - One byte opcode
  - Zero or more operands
    - Opcode indicates how many
- Can take operands from
  - Instruction
  - Current constant pool
  - Current frame local variables
  - Values on operand stack
    - Distinguish storage types and computation types



} index into  
constant  
pool or  
into local  
storage



# Implied Registers

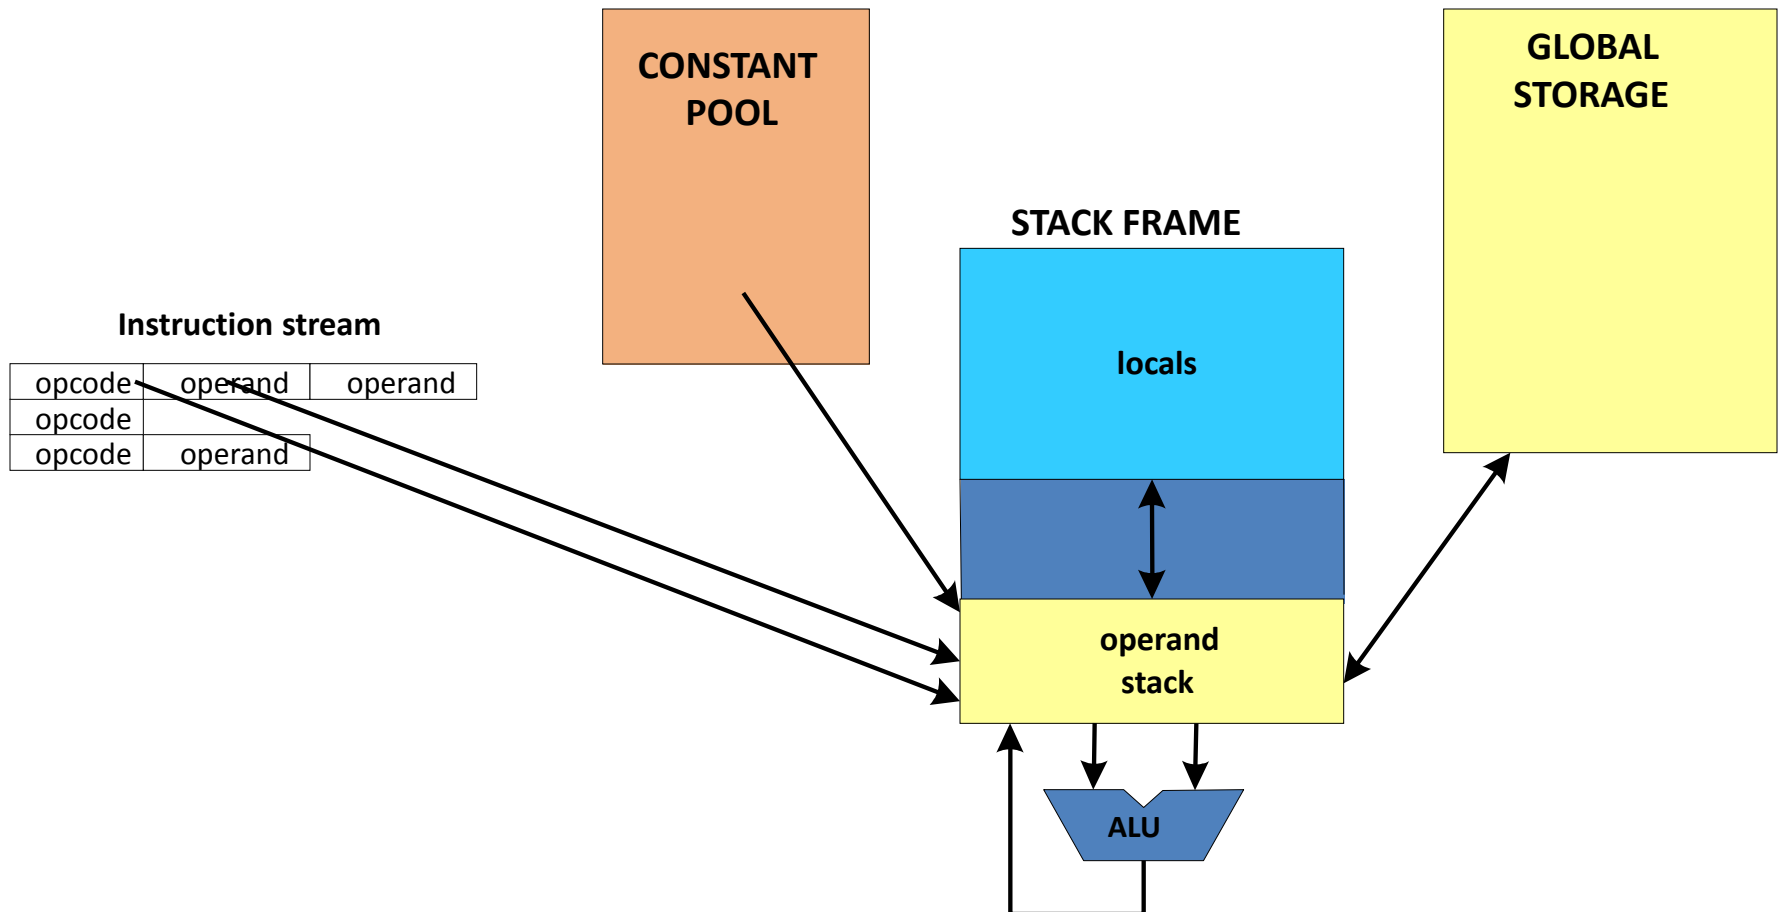
- Program Counter
- Local variable pointer
- Operand stack pointer
- Current frame pointer
- Constant pool base

# Instruction Types

- Pushing constants onto the stack
- Moving local variable contents to and from the stack
- Managing arrays
- Generic stack instructions (dup, swap, pop & nop)
- Arithmetic and logical instructions
- Conversion instructions
- Control transfer and function return
- Manipulating object fields
- Method invocation
- Miscellaneous operations
- Monitors

# Data Movement

- All data movement takes place through stack





# Bytecode Example

## PC instruction

```
0:  iconst_2 //pushes constant 2 onto operand stack
1:  aload_0  //pushes local variable 0 onto the stack
2:  getfield #2; //object ref on the stack, entry 2 on constant pool gives descr
5:  iconst_0
6:  iaload
7:  aload_0
8:  getfield #2;
11: iconst_1
12: iaload
13: iadd
14: imul
15: ireturn
```

# Stack Tracking

- Operand stack at any point in program have:

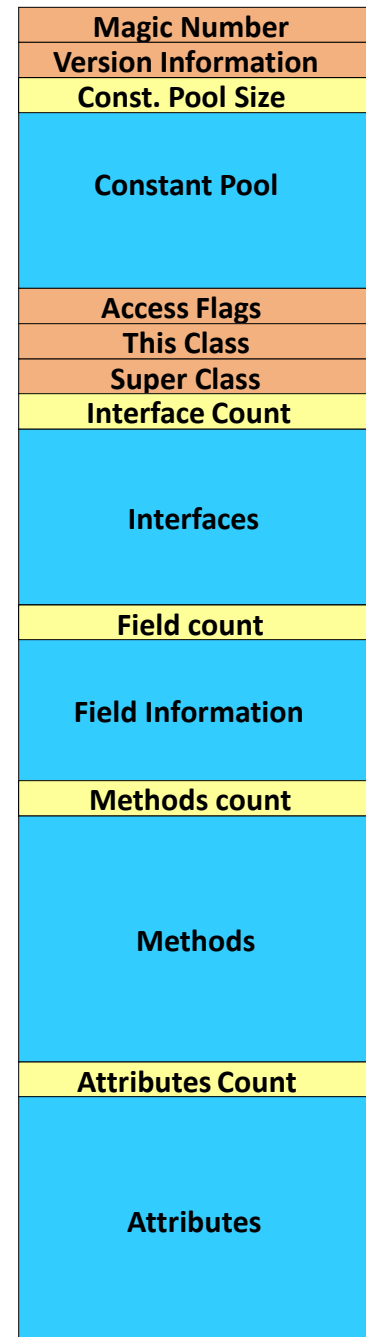
- Same number of operands
- Of same types
- In same order

*Regardless of control flow path getting there*

- Helps with static type checking by the loader

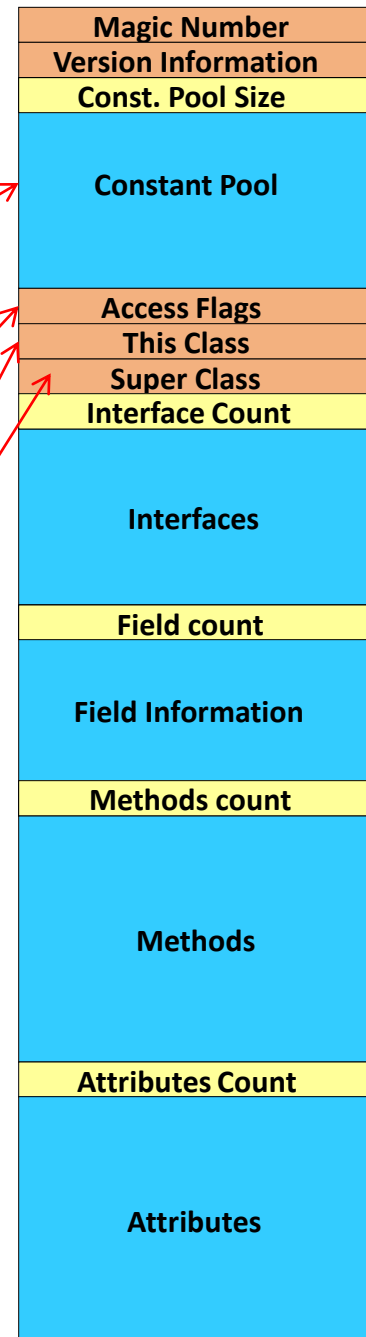
# Binary Classes

- Formal ISA Specification
- Magic number and header
- Major regions preceded by counts
  - Constant pool
  - Interfaces
  - Field information
  - Methods
  - Attributes



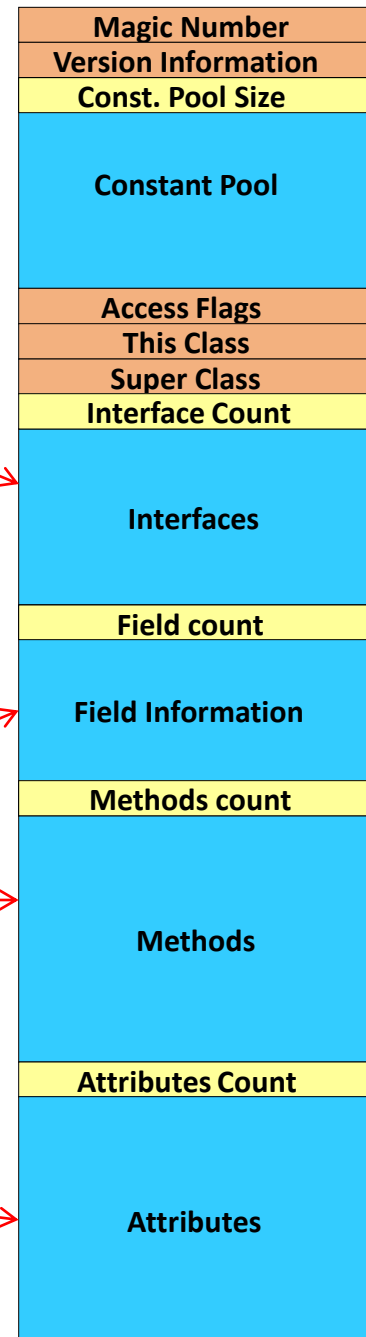
# Binary Classes

- Holds all constant values and references used by the methods that are to follow.
- Provides access information, example:
  - whether public
  - whether interface
  - ...
- Given as indices in the constant pool



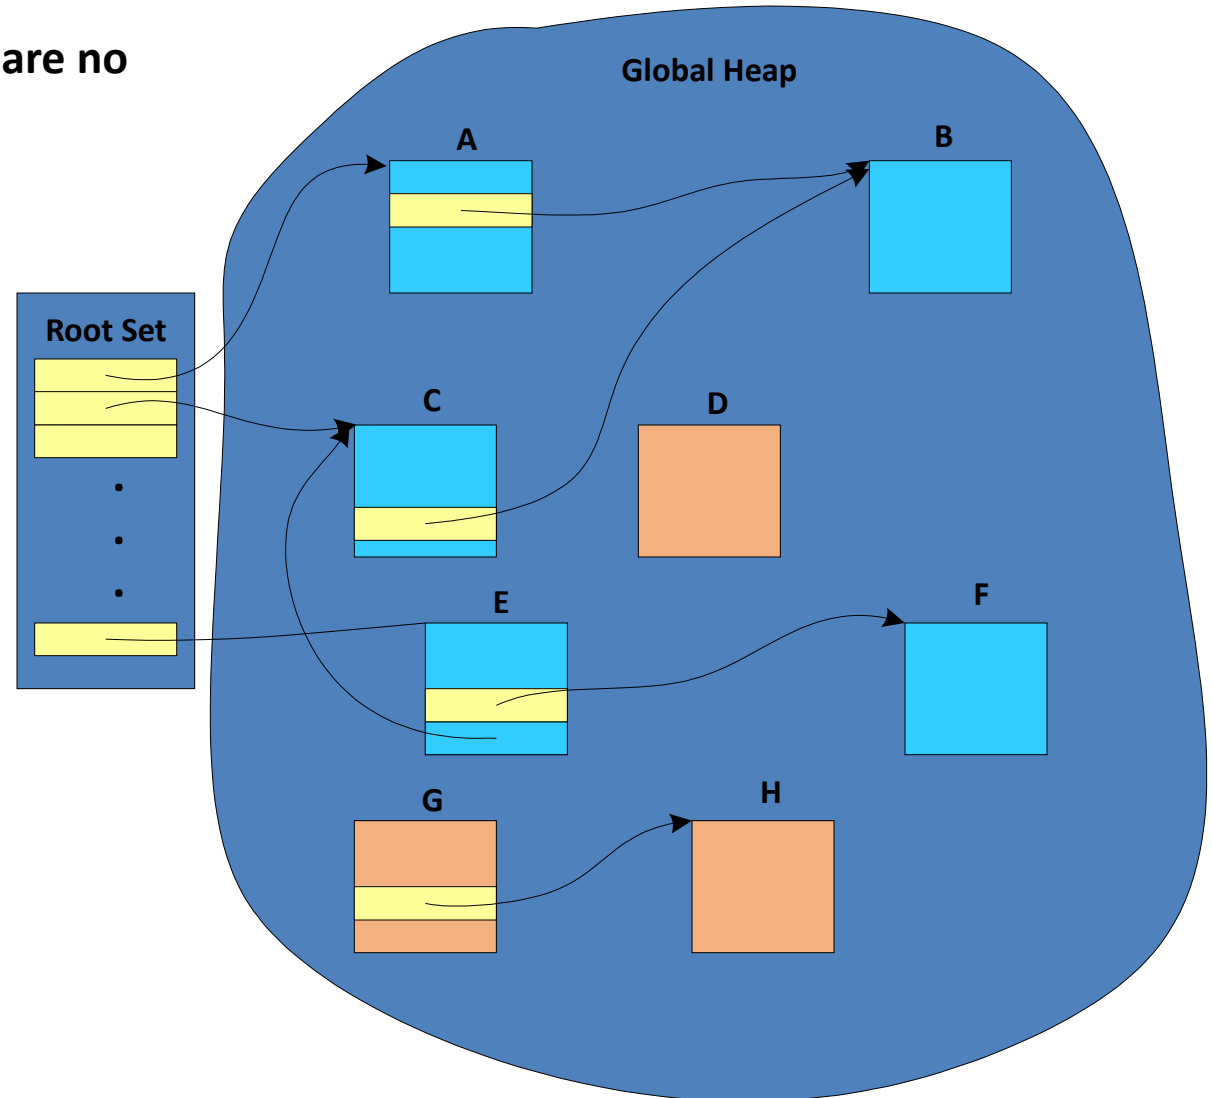
# Binary Classes

- Contains a number of references to the superinterfaces to this class
  - Given as indices in the constant pool
  - The constant pool entries are references to the interfaces
- Contains the specifications of the fields declared in this class
- The information regarding each method, as well as the methods themselves (encoded as bytecode)
- Contains detailed information regarding the previous sections



# A Note About: Garbage Collection

- ❑ Garbage: objects that are no longer accessible
- ❑ Examples:  
D, G, H

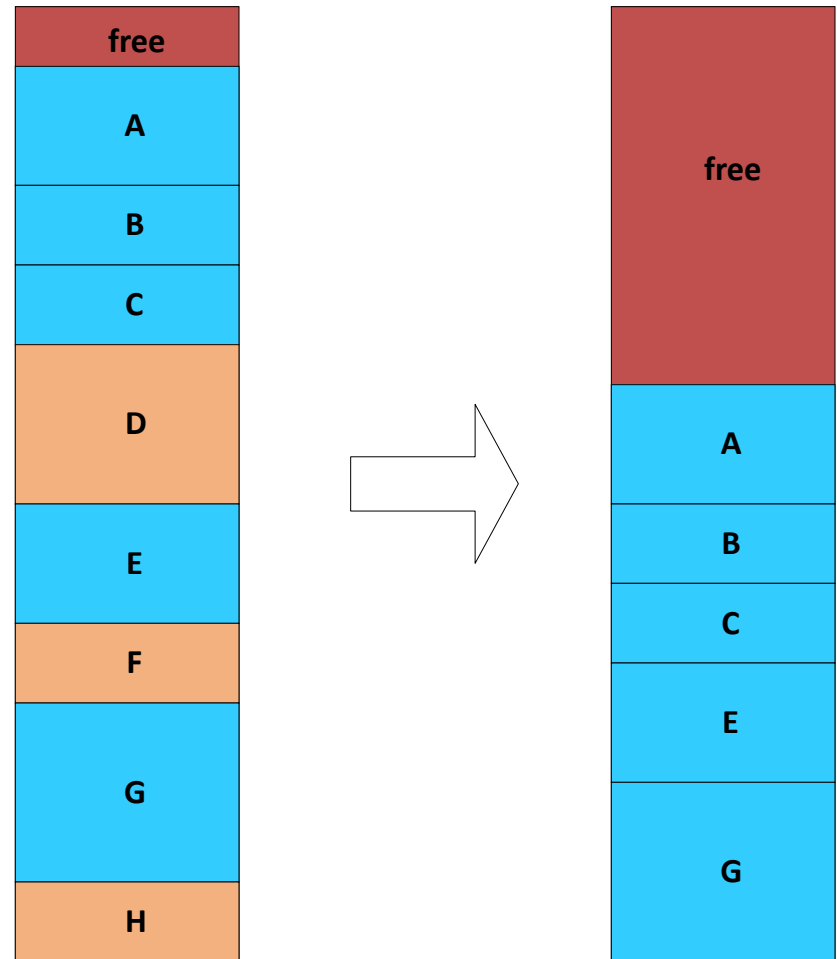


# Garbage Collection

- A large topic on its own
- Mark and sweep
  - Start with *root set* of references
    - On stack, static objects, constant pool
  - Trace and mark all reachable objects
- Sweep through heap, collecting marked objects
  - Keep free space in linked list
    - Advantage: Fast
  - Does not require moving object/pointers
- Disadvantage:
  - Discontiguous free space, fragmentation
  - Allocate new objects from best-fit free list

# Compacting Collector

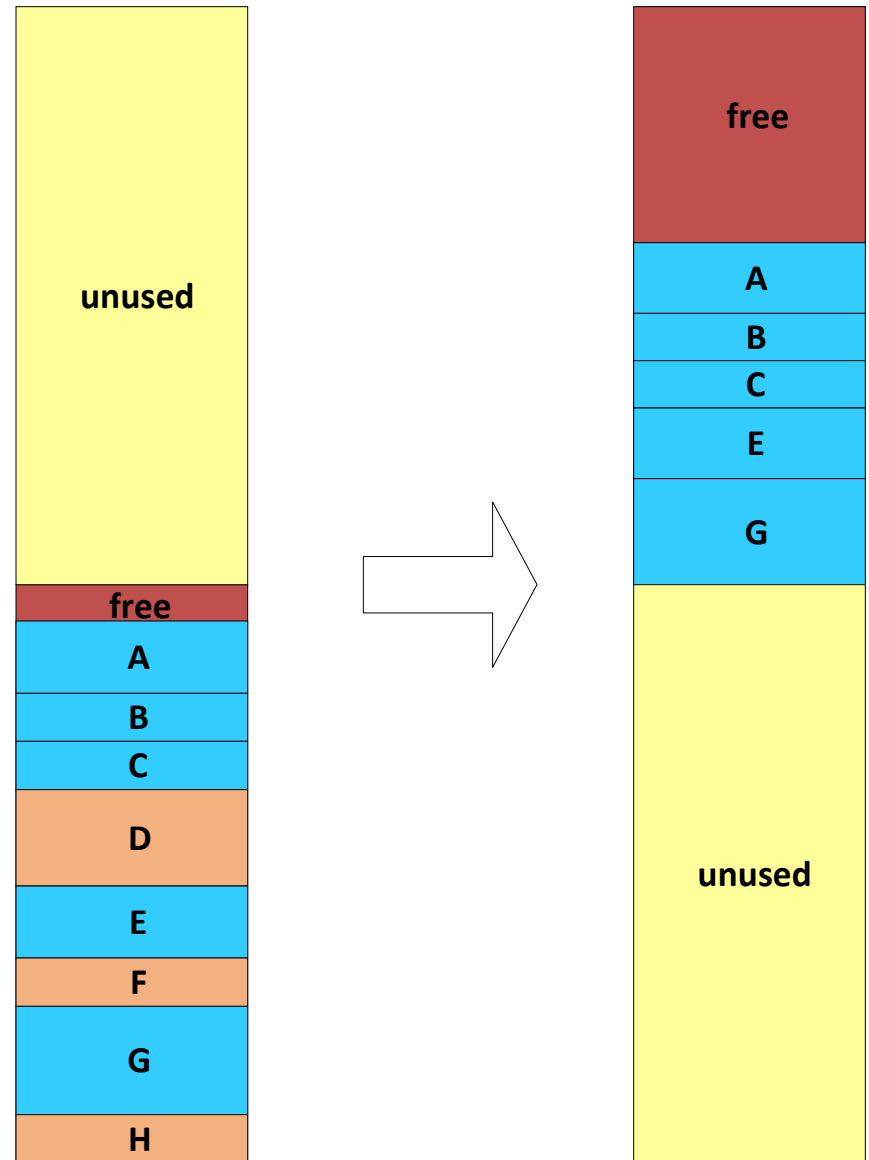
- Make free space contiguous
- Involves multiple passes through heap
- A lot of object movement => many pointer updates





# Copying Collector

- Divide heap into halves
- Collect when one half full
- Copy into unused half during sweep phase
- + Reduces passes through heap
- "Wastes" half of heap



# Generational Collector

- Divide heap into halves
  - “tenured” and “nursery”
- Collect nursery more frequently
- Move long-lived objects into tenured half
- Objects have either very long or very short lives

# JVM Bytecode Emulation

- Interpretation
  - Simple, fast startup, but slow
- Just-In-Time (JIT) Compilation
  - Compile each method when first touched
  - Simple, static optimizations
- Hot-Spot Compilation
  - Find frequently executed code
  - Apply more aggressive optimizations on that code
  - Typically phased with interpretation or JIT
- Dynamic Compilation
  - Based on Hot-Spot compilation
  - Use runtime information to optimize
  - More later...

# So JVM is:

- An abstract entity that gives meaning to class files
- Has many concrete implementations
  - Hardware
  - Interpreter
  - JIT compiler
- Persistence
  - An instance is created when an application starts
  - Terminates when the application finishes

# Conclusions

- HLL VM is built with portability as main goal:
  - Building a loader and JIT compiler is easier than building a full-fledged compiler
  - API and not ABI